

## C/UNIX Seminar - January 2022



# Day 3: I/O, signals and Introduction to IPC

Assistants Salles Machine 2024



## I/O in C

- High level buffered I/O

- Low-level I/O

- Working with files

## Inter Process Communication



# I/O in C



# I/O in C

- Two types on I/O: high and low level.
- Available in `<stdio.h>` or `<unistd.h>`.
- Concept of buffers for high level functions like *printf*.



# High level buffered I/O

The high level I/O is the most common one: *printf*, *puts*, etc... "Buffered" functions, meaning the content given to these functions is not always printed. The input is collected in buffers to write as much as possible at once. Then, everything gets outputted when:

- The buffer is full.
- A newline character is encountered.

There is actually a third one, but not relevant here.



# Formatted operations

- *printf(3)* and *scanf(3)* family.
- Variadic functions: can take an infinite amount of arguments.
- Format strings: *%d*, *%s*, ...

Feel free to look at their respective man page.



# Printing values

```
#include <stdio.h>
int main()
{
    int i = 42;
    float f = 3.141592;

    printf("i with a leading space = |%3d|\n", i);
    printf("i with a leading zero = |%03d|\n", i);
    printf("f = %g\n", f);
    fprintf(stderr, "Printing on stderr...\n");

    return 0;
}
```





# Reading values

```
#include <stdio.h>
int main()
{
    int age;

    puts("What is your age?");
    scanf("%d", &age);

    printf("You are %i years old.\n", age);

    return 0;
}
```



# Low-level I/O

The first notable difference with low-level I/O functions is that they require a *file descriptor*.

A file descriptor, or *fd* is a unique identifier - a number - for a file or other I/O resource, like network sockets or pipes.

The three FDs that interest us here are:

- 0: standard input *STDIN\_FILENO*.
- 1: standard output *STDOUT\_FILENO*.
- 2: standard error *STDERR\_FILENO*.



# Syscalls

*Syscalls* are "particular" functions. They are actually wrappers around existing OS operations and can be used as any other C function. Their man-page differ though as they are featured in section 2 of the manual, instead of the usual section 3.

Also, these functions are usually **blocking**, meaning that your program will wait for them to be finished before moving on.



# read(2)

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t count);
```

- Reads at most *count* bytes from *fd*. *count* should be lower or equal to the size of the buffer.
- Puts the read input in *buf*.
- Returns the number of bytes - characters in most cases - read. If it is equal to 0, reached *end-of-file*, everything was read; else if -1, an error occurred.



# write(2)

```
#include <unistd.h>
ssize_t write(int fd, void *buf, size_t count);
```

- Writes at most *count* bytes from *buf*. *count* should be lower or equal to the size of the buffer.
- Writes the content of the buffer in *fd*.
- Returns the number of bytes written. If it is equal to -1, an error occurred.



# Example

```
#include <err.h>
#include <unistd.h>

#define BUF_SIZE 128

int main()
{
    ssize_t r;
    char buf[BUF_SIZE];
    while ((r = read(STDIN_FILENO, buf, BUF_SIZE)) != 0)
    {
        if (r == -1)
            errx(3, "Reading failure.");
        r = write(STDOUT_FILENO, buf, r);
        if (r == -1)
            errx(3, "Writing failure.");
    }
    return 0;
}
```



# Error handling

As usual, you are expected to handle all possible errors. To do so, introducing: `<errno.h>` and its variable, *errno*.

A failing syscall will set *errno* to a specific value, and the possible errors are described in the different syscalls' man-pages. The explicit error can be obtained with *strerror(3)* or *err(3)*.



Also, there are two ways to notify errors:

- *err(3)* family: puts a message on stderr and exits the program.
- *warn(3)* family: puts a message on stderr but does not stop the program.





# Working with files

```
#include <fcntl.h>
int open(const char *path, int flags);
int open(const char *path, int flags, mode_t mode);
```

- Returns a *fd* to the given *path*.
- *flags* is used to specify the opening mode (reading, writing, etc)
- *mode* sets the permissions of a file when creating it.
- Returns -1 in case of an error.



```
#include <fcntl.h>

/* ... */
int fd;
if ((fd = open("foo.txt", O_RDONLY)) == -1)
    err(3, "error while opening %s", "foo.txt");
/* Do something with the content of the file */
close(fd);
```



# Getting file information: stat(2)

```
#include <sys/stat.h>
int stat(const char *path, struct stat*buf);
```

Allows to get:

- the size of a file,
- its permissions,
- the date of creation, modification, ...

Feel free to read the man-page to learn more about these functions and what they permit.



# Seeking in a file: seek(2)

```
#include <unistd.h>
off_t lseek(int fd, off_t offset, int whence);
```

In a sense, this function allows to move the *offset* of the resource designed by *fd*. An easier explanation would be like moving a cursor inside of a file.

*whence* specifies how the offset is to be interpreted:

- SEEK\_SET: file offset set to *offset* bytes.
- SEEK\_CUR: file offset set to its current location plus *offset*.
- SEEK\_END: file offset set to file size plus *offset*.



# Inter Process Communication



# Process

- An instance of a running program is called a process
- It can be identified by its PID (Process ID)
- Each instance of a running program has a unique PID



# Example

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(void)
{
    pid_t pid = getpid();

    printf("The pid of the current program is %d\n", pid);

    return 0;
}
```



# Or...

Even the terminal you use is a process that has it's own PID.  
You can access the PID using the \$\$ variable:

```
echo $$
```

(This will display the PID of the current terminal)





# IPC (Inter Process Communication)

Each process is independent which means each process has its own memory pool. That is why we need IPCs:

**Mechanisms** an operating system provides to allow the **processes to manage shared data** and "communicate" between each other



# Types of IPCs

1. Pipes
2. Shared Memory
3. Message Queue
4. Direct Communication (targeted)
5. Indirect communication (uses intermediary mailbox)
6. Message Passing (i.e.: sockets)
7. FIFO



# Signals

- Simple asynchronous Inter Process Communication (IPC)
- Designed mostly for exceptional behavior
- Most signals kill the process by default
- Most signals are catchable



# Signals

To send signal: `kill(2)` syscall (also `kill(1)`, a command!)

- Needs a signal number
- Needs a pid to send the signal to
- Requires ownership or enough privileges

Find the list of existing signal numbers with `man 7 signal` or using `'kill -l` on a terminal



# Basic Signal Handling

- Define functions to be called when signal arrives
- Not all signals can be handled



# Some Common Signals

- SIGHUP
- SIGINT
- **SIGKILL**
- **SIGSTOP**
- SIGTERM
- SIGSEGV
- SIGCHLD
- SIGCONT
- SIGBUS
- SIGILL



## More Flexible: sigaction(2)

- Provides more possibilities than signal(2)
- Provides finer control over signal handling
- Can give a lot more information to your handlers



